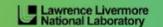
Automatic trace analysis with the Scalasca Trace Tools

The Scalasca Team
Jülich Supercomputing Centre







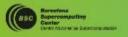










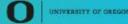












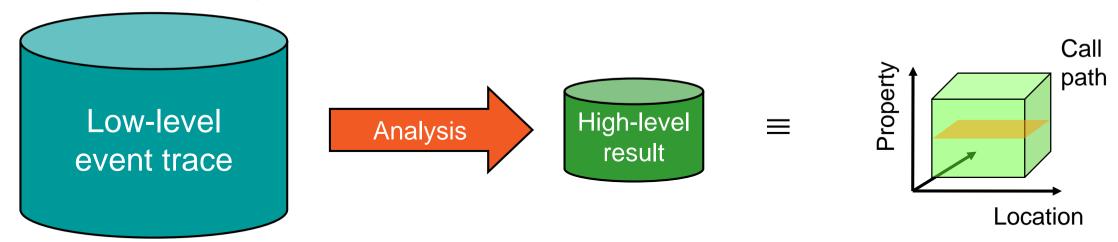




Automatic trace analysis

Idea

- Automatic search for patterns of inefficient behaviour
- Classification of behaviour & quantification of significance
- Identification of delays as root causes of inefficiencies



- Guaranteed to cover the entire event trace
- Quicker than manual/visual trace analysis
- Parallel replay analysis exploits available memory & processors to deliver scalability



Scalasca Trace Tools: Objective

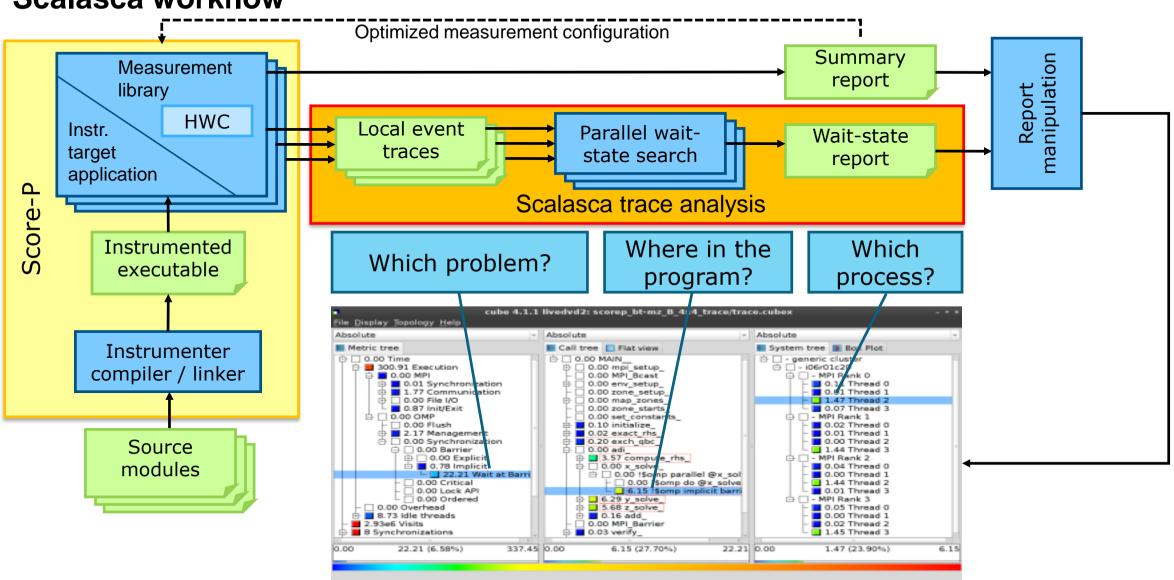
- Development of a scalable trace-based performance analysis toolset for the most popular parallel programming paradigms
 - Current focus: MPI, OpenMP, and (to a limited extend) POSIX threads
- Specifically targeting large-scale parallel applications
 - Demonstrated scalability up to 1.8 million parallel threads
 - Of course also works at small/medium scale
- Latest release:
 - Scalasca v2.6 coordinated with Score-P v7.0 (April 2021), also works with later versions

Scalasca Trace Tools: Features

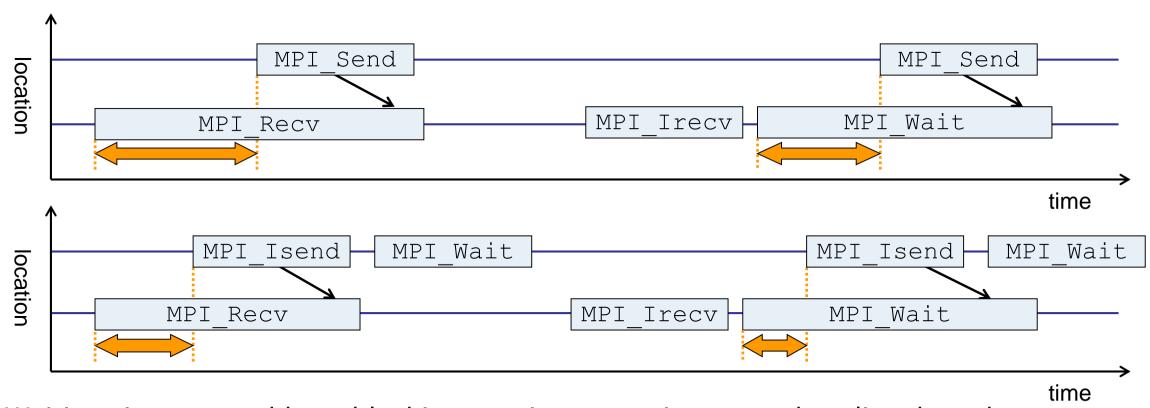
- Open source, 3-clause BSD license
- Fairly portable
 - HPC/Cray XT/XE/XK/XC/EX, IBM Blue Gene, SGI Altix, Fujitsu FX systems, Linux clusters (x86, Power, ARM), Intel Xeon Phi, ...
- Uses Score-P instrumenter & measurement libraries
 - Scalasca v2 core package focuses on trace-based analyses
 - Supports common data formats
 - Reads event traces in OTF2 format
 - Writes analysis reports in CUBE4 format
- Current limitations:
 - Unable to handle traces
 - with MPI thread level exceeding MPI_THREAD_FUNNELED
 - containing Memory events, CUDA/OpenCL device events (kernel, memcpy), SHMEM, or OpenMP nested parallelism
 - PAPI/rusage metrics for trace events are ignored



Scalasca workflow



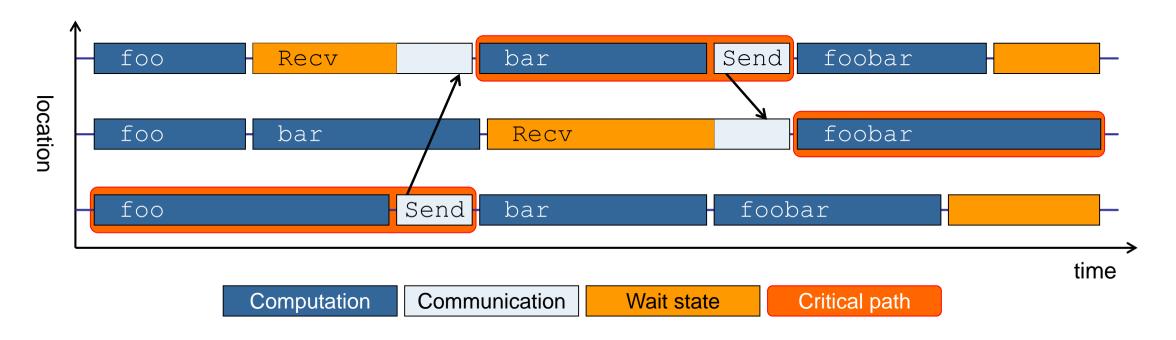
Example: "Late Sender" wait state



- Waiting time caused by a blocking receive operation posted earlier than the corresponding send
- Applies to blocking as well as non-blocking communication



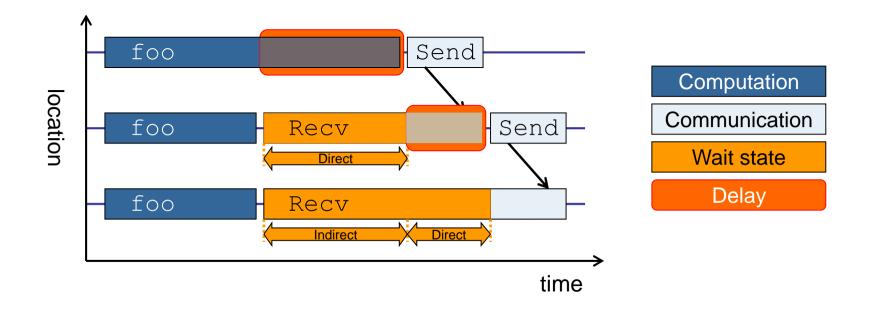
Example: Critical path



- Shows call paths and processes/threads that are responsible for the program's wall-clock runtime
- Identifies good optimization candidates and parallelization bottlenecks



Example: Root-cause analysis



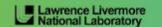
- Classifies wait states into direct and indirect (i.e., caused by other wait states)
- Identifies *delays* (excess computation/communication) as root causes of wait states
- Attributes wait states as delay costs

Hands-on: NPB-MZ-MPI / BT

































Scalasca command - One command for (almost) everything

```
% scalasca
Scalasca 2.6
Toolset for scalable performance analysis of large-scale parallel applications
usage: scalasca [OPTION]... ACTION <argument>...
    1. prepare application objects and executable for measurement:
       scalasca -instrument <compile-or-link-command> # skin (using scorep)
    2. run application under control of measurement system:
       scalasca -analyze <application-launch-command> # scan
    3. interactively explore measurement analysis report:
       scalasca -examine <experiment-archive | report > # square
Options:
  -c, --show-config
                         show configuration summary and exit
  -h, --help
                         show this help and exit
                         show actions without taking them
   -n, --dry-run
      --quickref
                         show quick reference quide and exit
      --remap-specfile show path to remapper specification file and exit
   -v, --verbose
                         enable verbose commentary
                         show version information and exit
   -V, --version
```

■ The `scalasca -instrument' command is deprecated and only provided for backwards compatibility with Scalasca 1.x., recommended: use Score-P instrumenter directly



Scalasca convenience command: scan / scalasca -analyze

```
% scan
Scalasca 2.6: measurement collection & analysis nexus
usage: scan {options} [launchcmd [launchargs]] target [targetargs]
     where {options} may include:
       Help
                 : show this brief usage message and exit.
 -v Verbose : increase verbosity.
 -n Preview : show command(s) to be launched but don't execute.
  -q Quiescent: execution with neither summarization nor tracing.
  -s Summary : enable runtime summarization. [Default]
 -t Tracing : enable trace collection and analysis.
       Analvze
                 : skip measurement to (re-)analyze an existing trace.
 -e exptdir
                 : Experiment archive to generate and/or analyze.
                   (overrides default experiment archive title)
 -f filtfile
                 : File specifying measurement filter.
                 : File that blocks start of measurement.
  -l lockfile
 -R #runs
                 : Specify the number of measurement runs per config.
 -M cfafile
                 : Specify a config file for a multi-run measurement.
```

Scalasca measurement collection & analysis nexus



Automatic measurement configuration

- scan configures Score-P measurement by automatically setting some environment variables and exporting them
 - E.g., experiment title, profiling/tracing mode, filter file, ...
 - Precedence order:
 - Command-line arguments
 - Environment variables already set
 - Automatically determined values
- Also, scan includes consistency checks and prevents corrupting existing experiment directories
- For tracing experiments, after trace collection completes then automatic parallel trace analysis is initiated
 - Uses identical launch configuration to that used for measurement (i.e., the same allocated compute resources)



Scalasca convenience command: square / scalasca -examine

```
% square
Scalasca 2.6: analysis report explorer
usage: square [OPTIONS] <experiment archive | cube file>
   -c <none | quick | full> : Level of sanity checks for newly created reports
                            : Force remapping of already existing reports
   -F
  -f filtfile
                            : Use specified filter file when doing scoring (-s)
                            : Skip display and output textual score report
  -s
                            : Enable verbose mode
                            : Do not include idle thread metric
   -n
                            : Aggregation method for summarization results of
   -S <mean | merge>
                              each configuration (default: merge)
   -T <mean | merge>
                            : Aggregation method for trace analysis results of
                              each configuration (default: merge)
                            : Post-process every step of a multi-run experiment
   -A
```

Scalasca analysis report explorer (Cube)



Recap: Local installation (Archer2)

- Setup access to Scalasca and associated tools
 - Source module use /work/y23/shared/tutorial/modulefiles to prepare the environment
 - Required for each shell session
 - Score-P and Scalasca installations are toolchain specific

```
% module restore PrgEnv-gnu
% module use /work/y23/shared/tutorial/modulefiles
% module load scalasca
```

- Check module avail scalasca for alternate Score-P/Scalasca modules available
- Copy tutorial sources to your personal workspace

```
% cd $WORK
% tar zxvf /work/y23/shared/tutorial/NPB3.3-MZ-MPI.tar.gz
% cd NPB3.3-MZ-MPI
```



BT-MZ summary measurement collection...

```
% cd bin.scorep
% cp ../jobscript/archer2/scalasca.sbatch .
% cat scalasca.sbatch
# Scalasca nexus configuration for profiling
#NEXUS="scalasca -analyze"
# Scalasca nexus configuration for profiling
#NEXUS="scalasca -analyze -t"
# Score-P measurement configuration
export SCOREP FILTERING FILE=../config/scorep.filt
#export SCOREP TOTAL MEMORY=100M
# run the application
scalasca -analyze srun ./bt-mz C.8
```

Change to
 directory with the
 Score-P
 instrumented
 executable and
 edit the job script

```
Hint:
scan = scalasca -analyze
-s = profile/summary (def)
```

Submit the job



BT-MZ summary measurement

```
S=C=A=N: Scalasca 2.6 runtime summarization
S=C=A=N: ./scorep bt-mz C 8x6 sum experiment archive
S=C=A=N: Thu Jun 10 11:48:50 2021: Collect start
srun ./bt-mz C 8x6
NAS Parallel Benchmarks (NPB3.3-MZ-MPI) -
   BT-MZ MPI+OpenMP Benchmark
Number of zones: 8 \times 8
Iterations: 200 dt: 0.000100
Number of active processes:
 [... More application output ...]
S=C=A=N: Thu Jun 10 11:49:02 2021: Collect done (status=0) 12s
S=C=A=N: ./scorep bt-mz C 8x6 sum complete.
```

- Run the application using the Scalasca measurement collection & analysis nexus prefixed to launch command
- Creates experiment directory:scorep_bt-mz_C_8x6_sum

BT-MZ summary analysis report examination

Score summary analysis report

```
% square -s scorep_bt-mz_C_8x6_sum
INFO: Post-processing runtime summarization report (profile.cubex)...
INFO: Score report written to ./scorep_bt-mz_C_8x6_sum/scorep.score
```

Post-processing and interactive exploration with Cube

```
% square scorep_bt-mz_C_8x6_sum
INFO: Displaying ./scorep_bt-mz_C_8x6_sum/summary.cubex...

[GUI showing summary analysis report]
```

Hint:

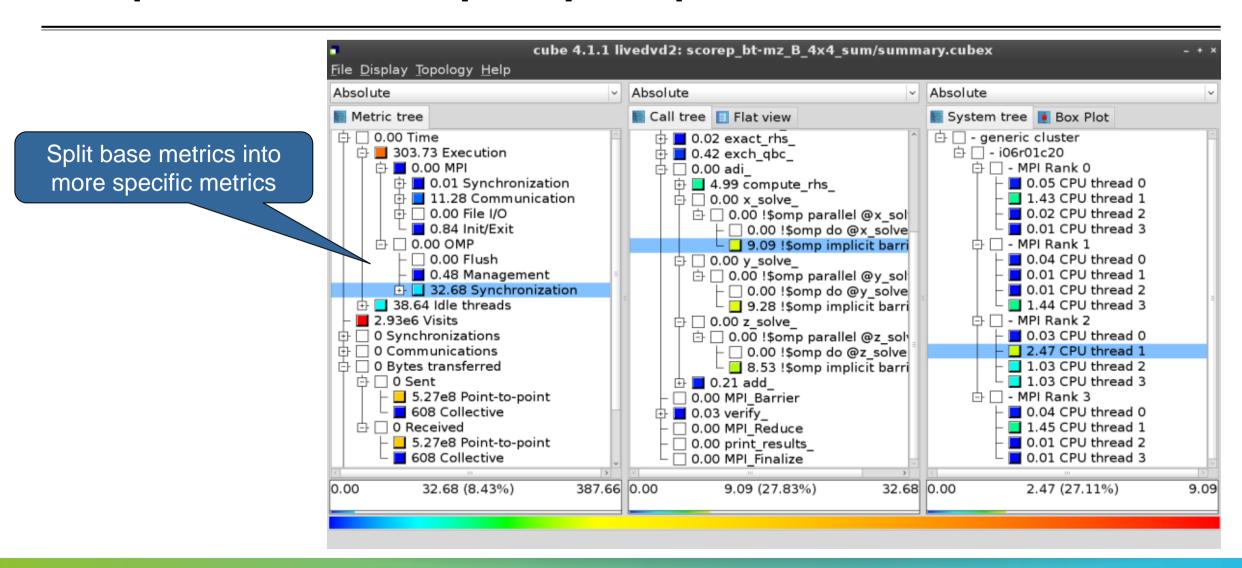
Copy 'summary.cubex' to local system (laptop) using 'scp' to improve responsiveness of GUI

 The post-processing derives additional metrics and generates a structured metric hierarchy

21



Post-processed summary analysis report



Performance analysis steps

- 0.0 Reference preparation for validation
- 1.0 Program instrumentation
- 1.1 Summary measurement collection
- 1.2 Summary analysis report examination
- 2.0 Summary experiment scoring
- 2.1 Summary measurement collection with filtering
- 2.2 Filtered summary analysis report examination
- 3.0 Event trace collection
- 3.1 Event trace examination & analysis



BT-MZ trace measurement collection...

```
% cd bin.scorep
% cp ../jobscript/archer2/scalasca.sbatch .
% vim scalasca.sbatch
# Scalasca nexus configuration for profiling
#NEXUS="scalasca -analyze"
# Scalasca nexus configuration for profiling
#NEXUS="scalasca -analyze -t"
# Score-P measurement configuration
export SCOREP FILTERING FILE=../config/scorep.filt
export SCOREP TOTAL MEMORY=100M
# run the application
scalasca -analyze -t srun ./bt-mz C.8
```

 Change to directory with the Score-P instrumented executable and edit the job script

- Add "-t" to the scan command
- Submit the job

% sbatch -reservation=ta031_192 scalasca.sbatch

BT-MZ trace measurement ... collection

```
S=C=A=N: Scalasca 2.6 trace collection and analysis S=C=A=N: ./scorep_bt-mz_C 8x6 trace experiment archive S=C=A=N: Thu Jun 10 12:05:30 2021: Collect start srun ./bt-mz_C.8

NAS Parallel Benchmarks (NPB3.3-MZ-MPI) - BT-MZ MPI+OpenMP \>Benchmark

Number of zones: 16 x 16
Iterations: 200 dt: 0.000100
Number of active processes: 8

[... More application output ...]

S=C=A=N: Thu Jun 10 12:05:44 2021: Collect done (status=0) 14s
```

- Using new experiment directory
 scorep bt-mz C 8x6 trace
- Starts measurement with collection of trace files...

BT-MZ trace measurement ... analysis

```
S=C=A=N: Thu Jun 10 12:05:44 2021: Analyze start
srun scout.hyb --time-correct \
        ./scorep bt-mz C 8x6 trace/traces.otf2
SCOUT (Scalasca 2.6)
Analyzing experiment archive ./scorep bt-mz C 8x6 trace/traces.otf2
Opening experiment archive ... done (0.002s).
Reading definition data
Reading event trace data
Preprocessing
Timestamp correction
Analyzing trace data
Writing analysis report

... done (0.002s).
done (0.113s).
... done (0.179s).
... done (0.431s).
... done (5.174s).
                                       : 422.312MB
Max. memory usage
            # passes : 1
# violated : 0
Total processing time : 6.140s
S=C=A=\bar{N}: Thu Jun 10 12:05:51 2021: Analyze done (status=0) 7s
```

 Continues with automatic (parallel) analysis of trace files



BT-MZ trace analysis report exploration

 Produces trace analysis report in the experiment directory containing trace-based wait-state metrics

```
% square scorep_bt-mz_C_8x6_trace
INFO: Post-processing runtime summarization report (profile.cubex)...
INFO: Post-processing trace analysis report (scout.cubex)...
INFO: Displaying ./scorep_bt-mz_C_8x6_trace/trace.cubex...

[GUI showing trace analysis report]
```

Hint:

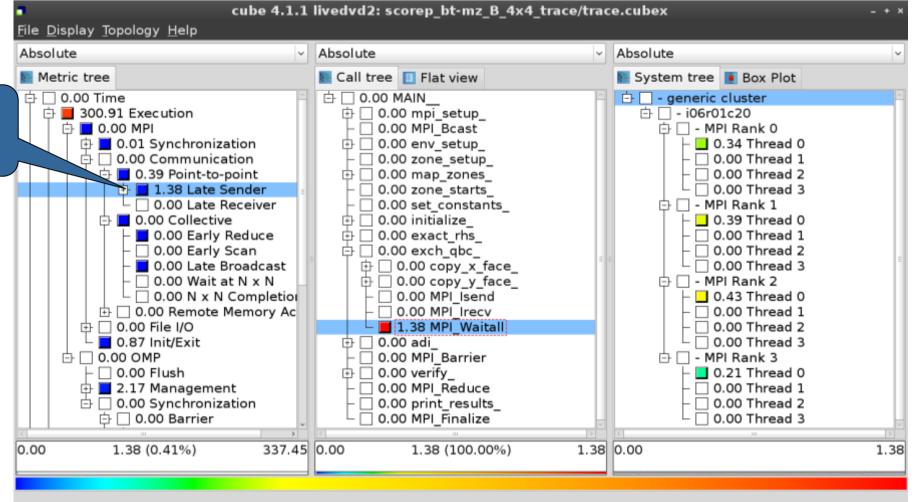
Run 'square -s' first and then copy 'trace.cubex' to local system (laptop) using 'scp' to improve responsiveness of GUI



Post-processed trace analysis report



Additional trace-based metrics in metric hierarchy

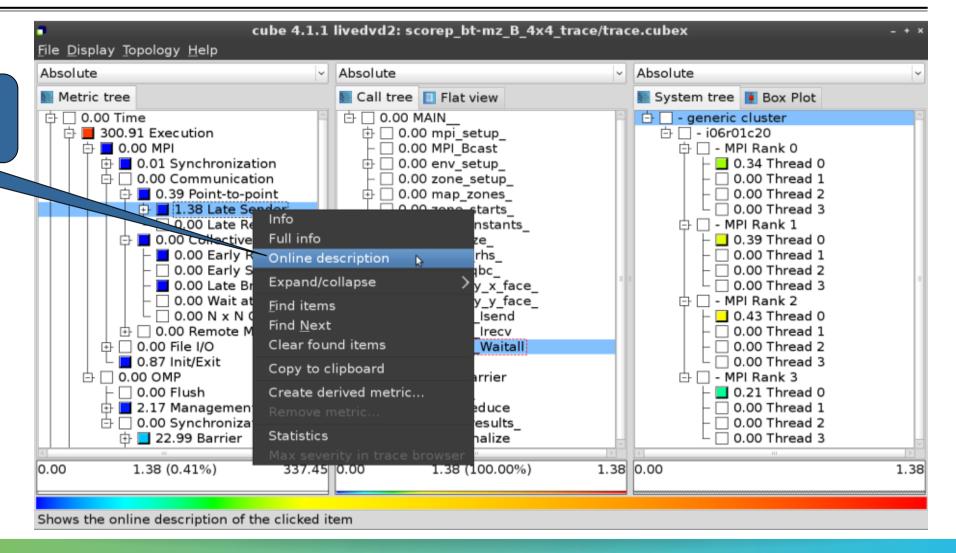




Online metric description



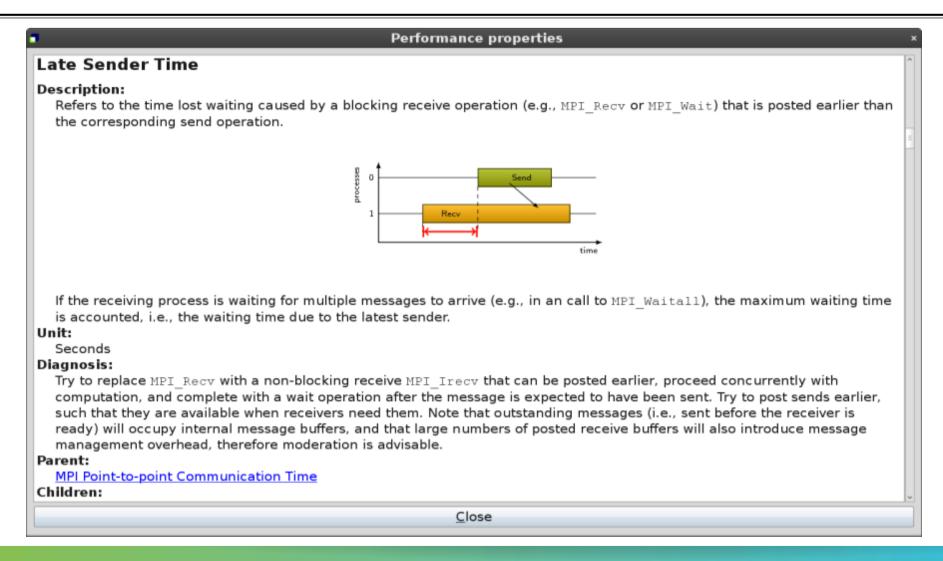
Access online metric description via context menu





Online metric description

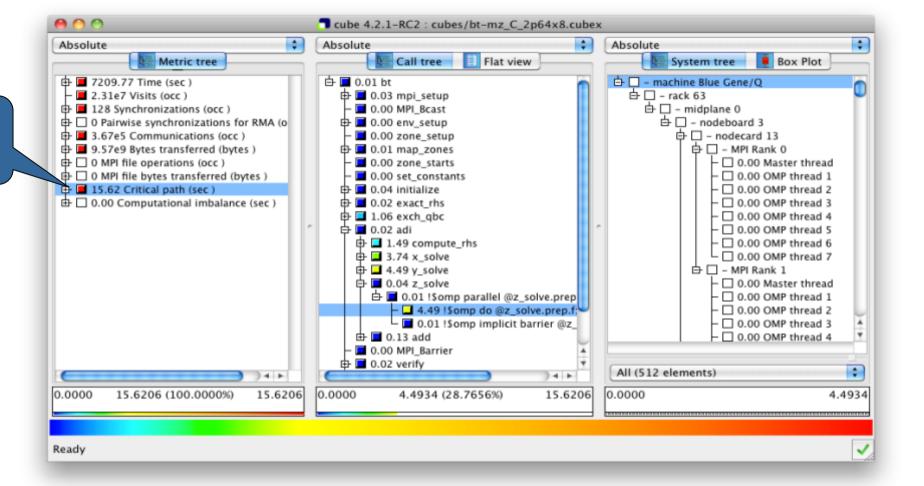




Critical-path analysis



Critical-path profile shows wall-clock time impact

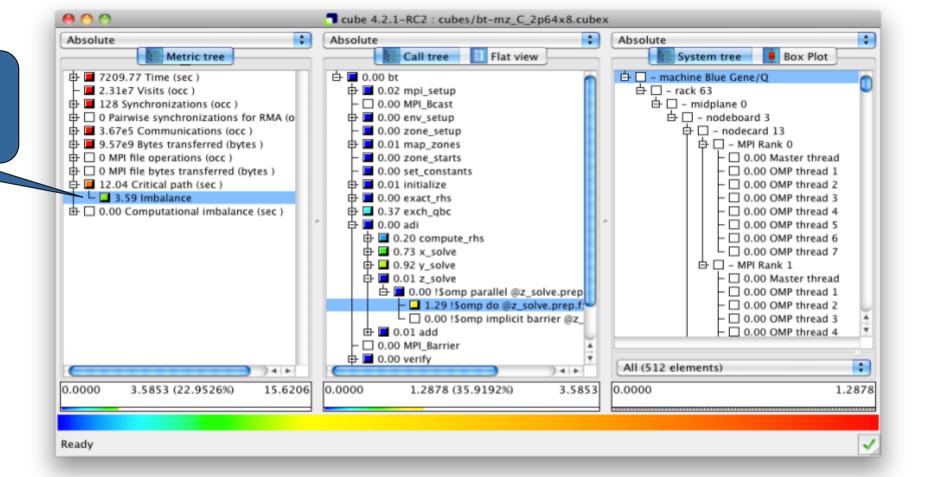




Critical-path analysis

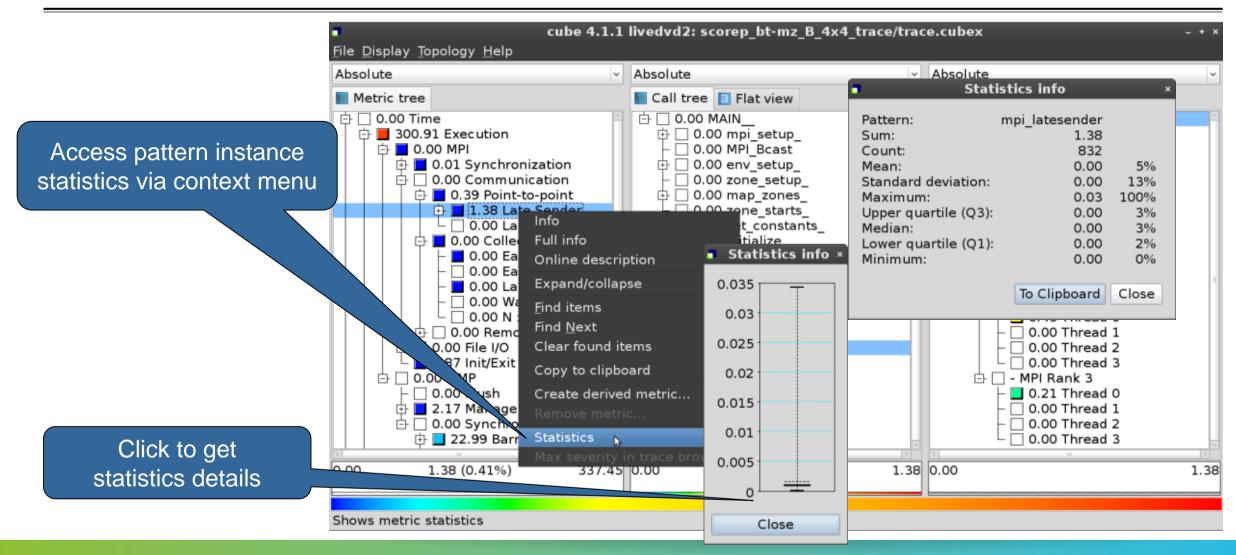


Critical-path imbalance highlights inefficient parallelism



Pattern instance statistics



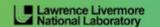


Demo: TeaLeaf case study



























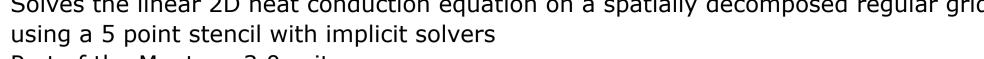






Case study: TeaLeaf

- HPC mini-app developed by the UK Mini-App Consortium
 - Solves the linear 2D heat conduction equation on a spatially decomposed regular grid using a 5 point stencil with implicit solvers





- Part of the Mantevo 3.0 suite
- Available on GitHub: https://uk-mac.github.io/TeaLeaf/
- Measurements of TeaLeaf reference v1.0 taken on Jureca cluster @ JSC
 - Using Intel 19.0.3 compilers, Intel MPI 2019.3, Score-P 5.0, and Scalasca 2.5
 - Run configuration
 - 8 MPI ranks with 12 OpenMP threads each
 - Distributed across 4 compute nodes (2 ranks per node)
 - Test problem "5": 4000 × 4000 cells, CG solver

```
% cp /home/vihps/public/case studies/♥
                              scorep_tea_leaf baseline 8x12 trace .
 cube scorep tea leaf baseline 8x12 trace/trace.cubex
                           [GUI showing post-processed trace analysis report]
```

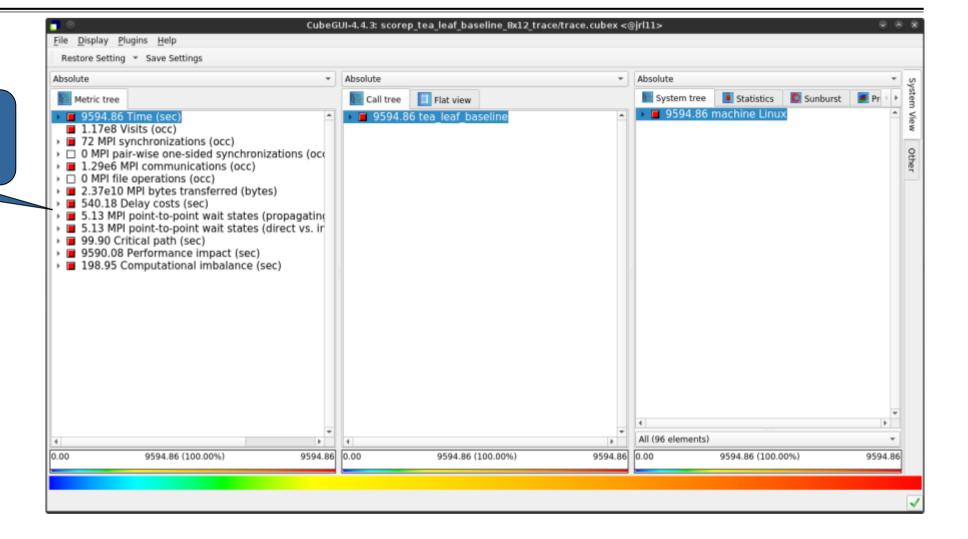
Hint:

Copy 'trace.cubex' to local system (laptop) using 'scp' to improve responsiveness of GUI

Scalasca analysis report exploration (opening view)



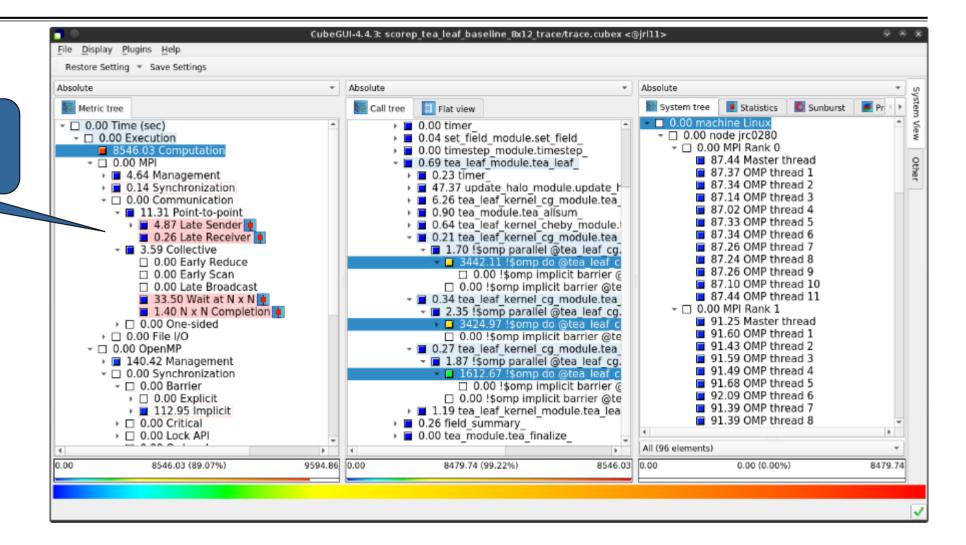
Additional top-level metrics produced by the trace analysis...



Scalasca wait-state metrics



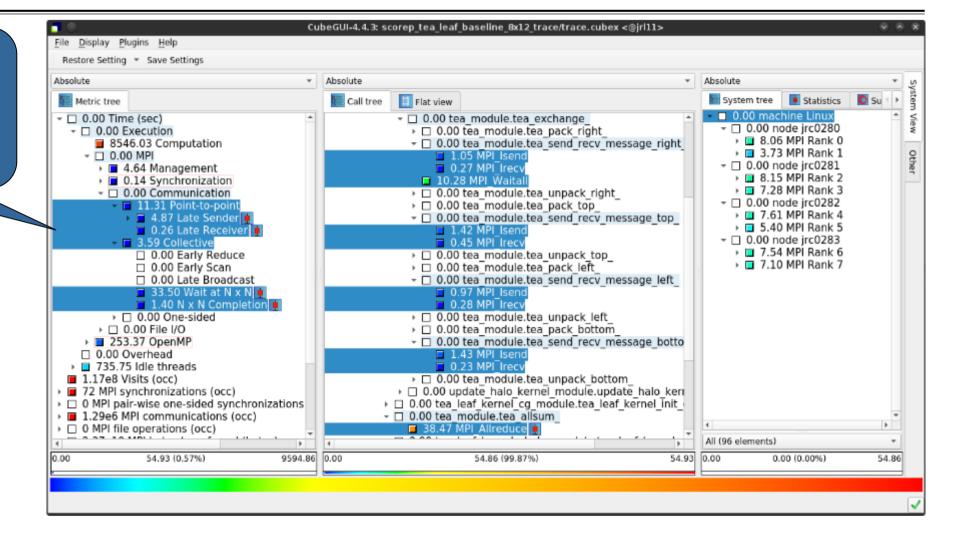
...plus additional waitstate metrics as part of the "Time" hierarchy



TeaLeaf Scalasca report analysis (I)



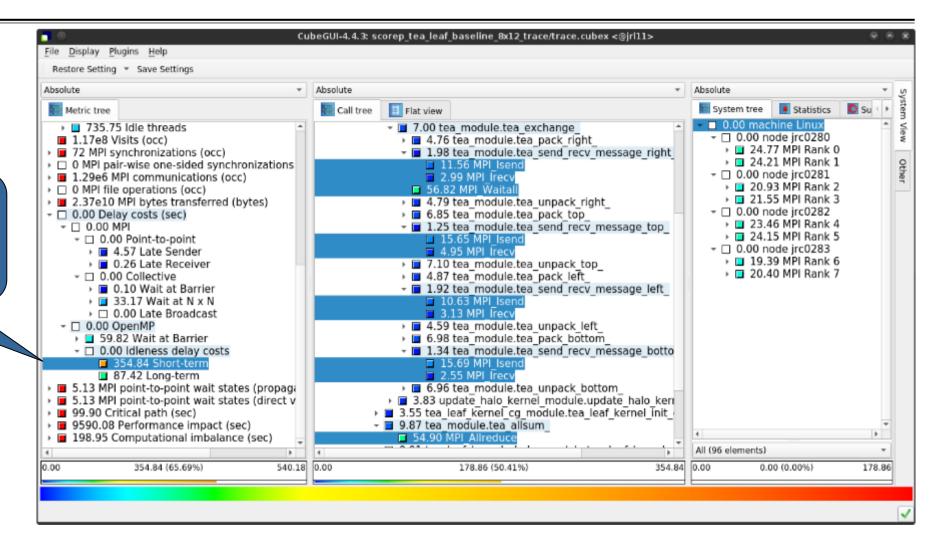
While MPI communication time and wait states are small (~0.6% of the total execution time)...



TeaLeaf Scalasca report analysis (II)



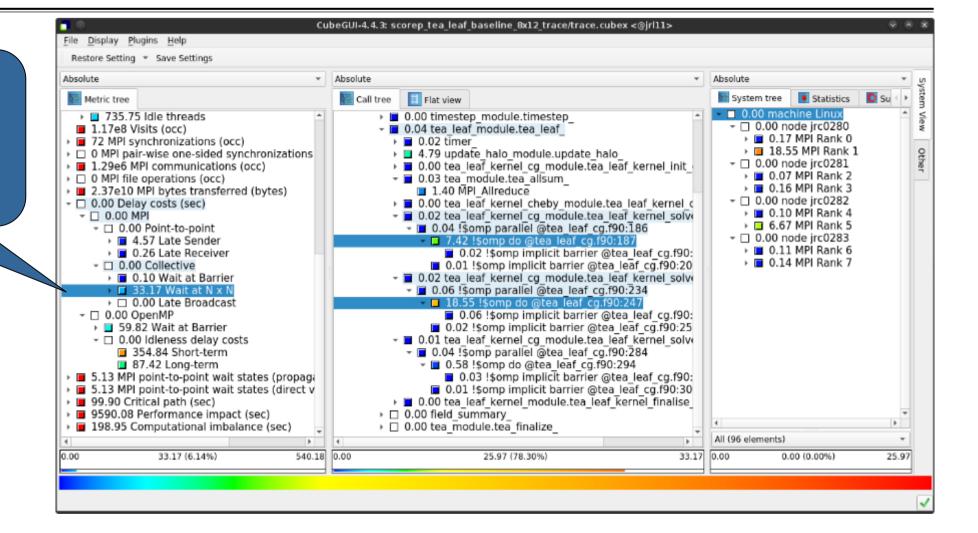
...they directly cause a significant amount of the OpenMP thread idleness



TeaLeaf Scalasca report analysis (III)



The "Wait at NxN" collective wait states are mostly caused by the first 2 OpenMP do loops of the solver (on ranks 5 & 1, resp.)...

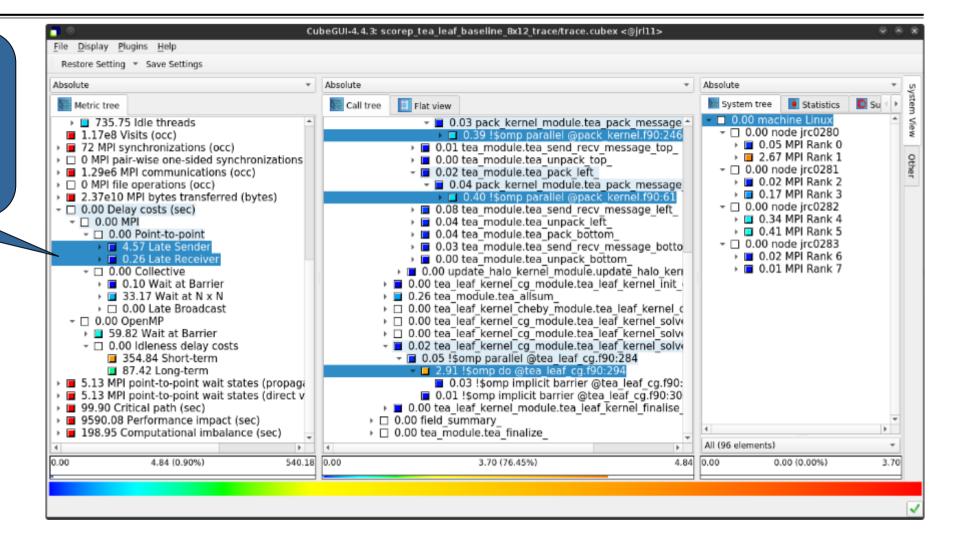




TeaLeaf Scalasca report analysis (IV)



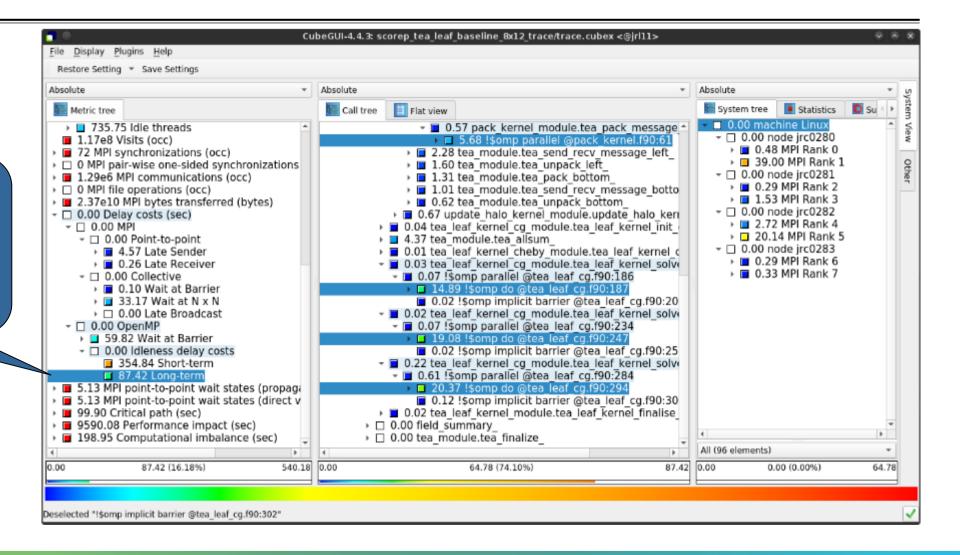
...while the MPI pointto-point wait states are caused by the 3rd solver do loop (on rank 1) and two loops in the halo exchange



TeaLeaf Scalasca report analysis (V)



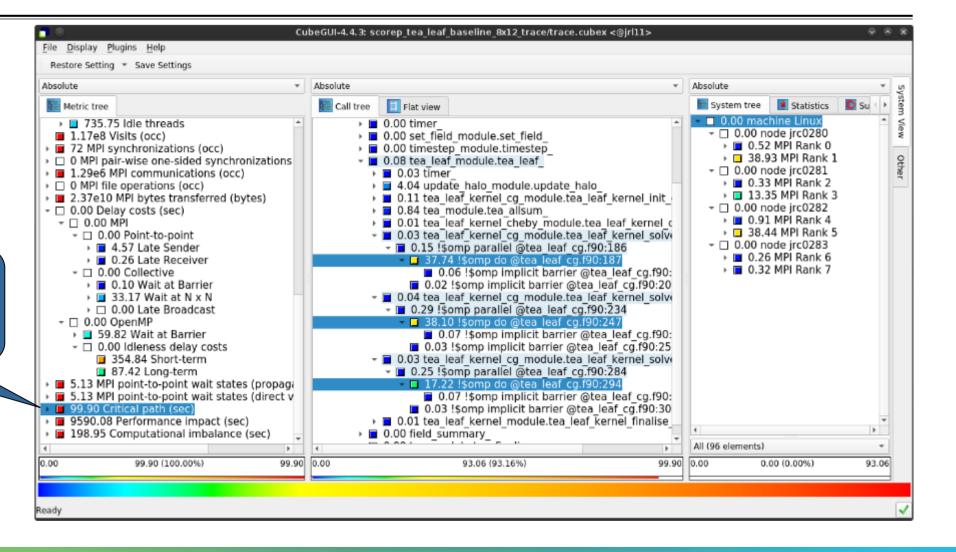
Various OpenMP do loops (incl. the solver loops) also cause OpenMP thread idleness on other ranks via propagation



TeaLeaf Scalasca report analysis (VI)



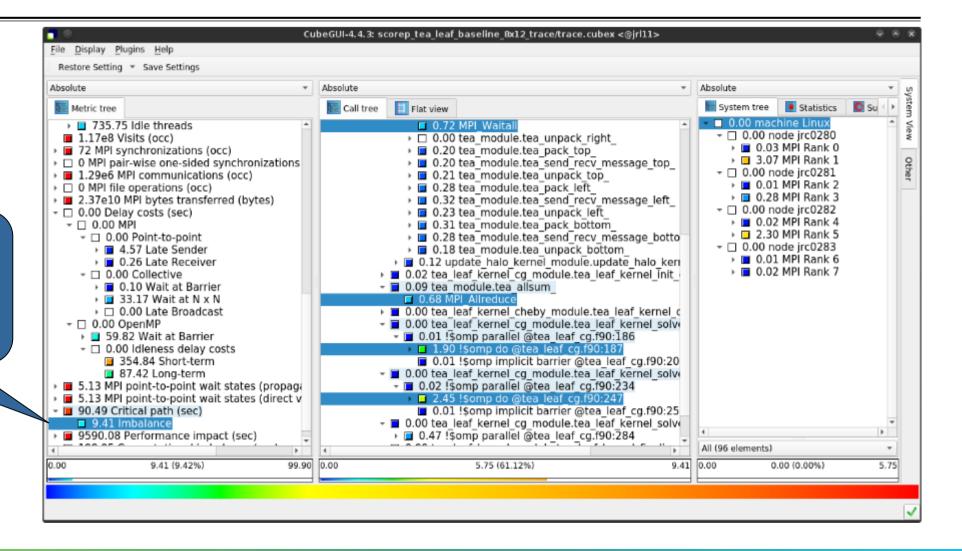
The Critical Path also highlights the three solver loops...



TeaLeaf Scalasca report analysis (VII)



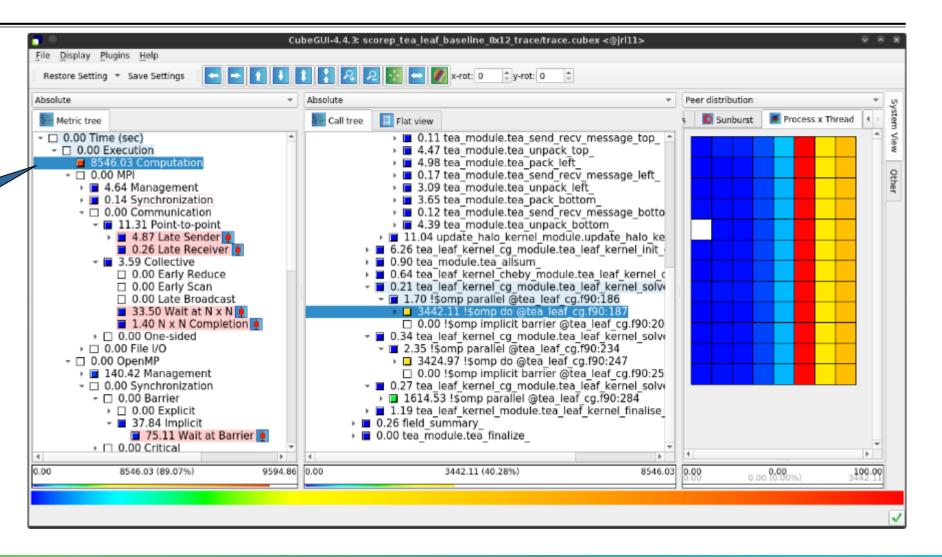
...with imbalance (time on critical path above average) mostly in the first two loops and MPI communication



TeaLeaf Scalasca report analysis (VIII)



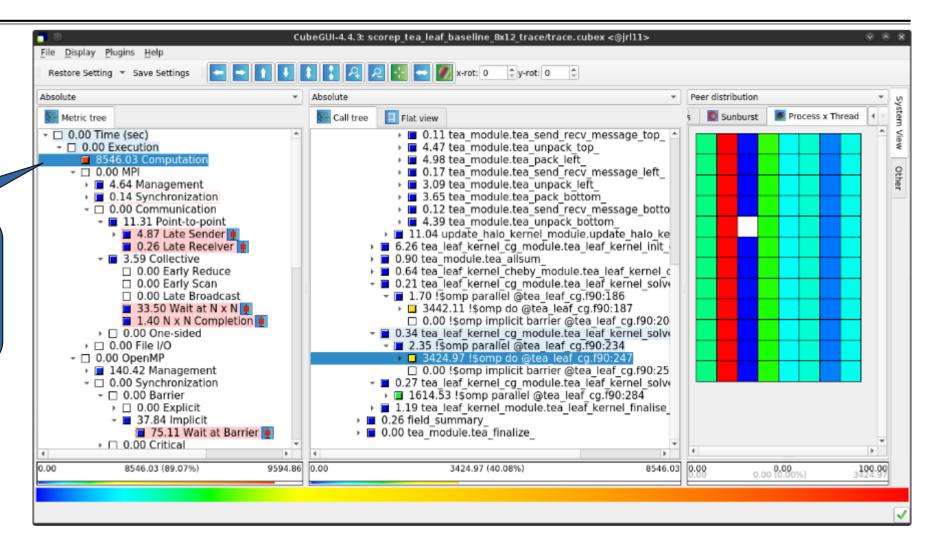
Computation time of 1st...



TeaLeaf Scalasca report analysis (IX)



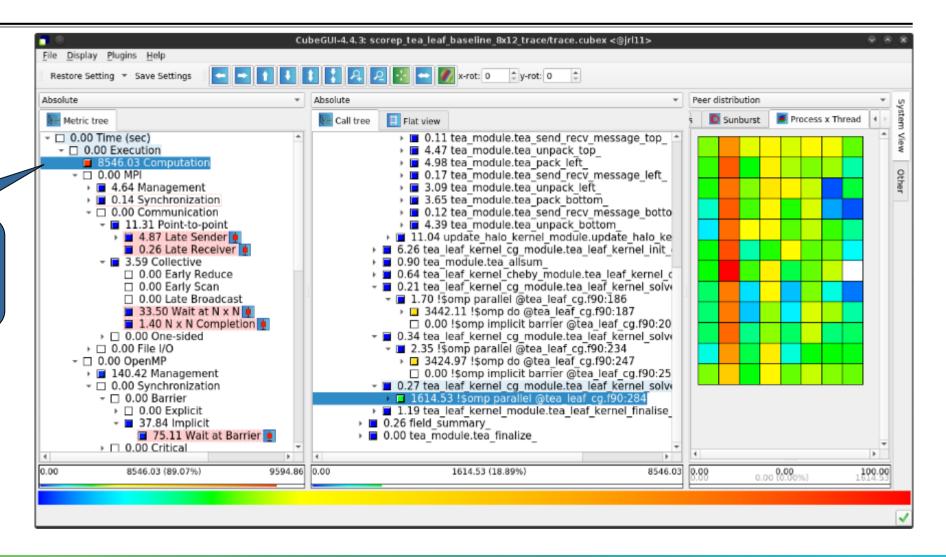
...and 2nd do loop mostly balanced within each rank, but vary considerably across ranks...



TeaLeaf Scalasca report analysis (X)



...while the 3rd do loop also shows imbalance within each rank



TeaLeaf analysis summary

- The first two OpenMP do loops of the solver are well balanced within a rank, but are imbalanced across ranks
 - → Requires a global load balancing strategy
- The third OpenMP do loop, however, is imbalanced within ranks,
 - causing direct "Wait at OpenMP Barrier" wait states,
 - which cause indirect MPI point-to-point wait states,
 - which in turn cause OpenMP thread idleness
 - → Low-hanging fruit
- Adding a SCHEDULE (quided) clause reduced
 - the MPI point-to-point wait states by ~66%
 - the MPI collective wait states by ~50%
 - the OpenMP "Wait at Barrier" wait states by ~55%
 - the OpenMP thread idleness by ~11%
 - → Overall runtime (wall-clock) reduction by ~5%



Scalasca Trace Tools: Further information

- Collection of trace-based performance tools
 - Specifically designed for large-scale systems
 - Features an automatic trace analyzer providing wait-state, critical-path, and delay analysis
 - Supports MPI, OpenMP, POSIX threads, and hybrid MPI+OpenMP/Pthreads
- Available under 3-clause BSD open-source license
- Documentation & sources:
 - https://www.scalasca.org
- Contact:
 - mailto: scalasca@fz-juelich.de

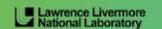


Exercises (if you don't have your own code)



































Warm-up

- Build the BT-MZ example code for class (i.e., problem size) "D"
 - Perform a baseline measurement w/o instrumentation (should run in ~190s)
 - Re-build the executable with Score-P instrumentation
- Repeat the hands-on exercise with the new executable
 - Perform a summary measurement
 - Score the summary measurement result
 - Adjust the measurement configuration appropriately
 - Perform a trace measurement and analysis



Trace analysis report examination

- What is the poportion of computation time vs. parallelization overheads?
- Which code regions are mostly responsible for the overall execution time?
- Are there any load balancing issues?
- If so, in which routines?
- What are the most significant wait states/parallelization overheads?
- What are their root causes?



Optimization

- What are possible optimizations?
 - Hint: Take a look at the TeaLeaf case study
- Modify the source code to apply those and re-do the measurement
 - Don't worry it's straightforward even if you don't know the code ;-)
 - Remember: One step at a time!
- How did the performance change?
 - The cube_diff tool or the "Cube Diff" plugin of the GUI (see the File → Context-free plugin menu) may come in handy here